

Study title

From public record to portfolio intelligence: validating LLM classification of the UK's register of accredited research projects

Research team

Dr Balint Stewart, Researcher — balint.stewart@ucl.ac.uk

Dr Joseph Lam, Principal Investigator — joseph.lam@ucl.ac.uk

UCL Great Ormond Street Institute of Child Health, Faculty of Population Health Sciences, University College London

UCL ethics application

Project ID: 5004

You are being invited to take part in a research study. Before you decide, please read this sheet so you understand what taking part involves. You may contact us with any questions before deciding.

Participant information

What is the study about?

We have developed a public dashboard based on the UK Statistics Authority's register of research projects accredited under the Digital Economy Act 2017. The dashboard makes this research easier to search and understand and includes experimental classifications of projects by research domain and analytical purpose.

The classifications are generated by a large language model (LLM) using only the information published in the public register entry. They are experimental and are not presented as authoritative descriptions of any project. This validation study assesses how accurately those classifications represent the projects, and whether apparent problems arise from classification error, limited information in the public register entry, or gaps in the classification framework itself.

Dashboard: <https://adr-dea-project.onrender.com/>

Why have I been invited?

You are publicly named as a researcher on one or more eligible projects in the register. Researchers directly involved in a project can provide context that may not be visible from the public register entry alone.

What would taking part involve?

You will receive one personalised link. After confirming that you are the intended recipient, reading the Participant Information Sheet and agreeing to take part, you will see a list of the projects on which you are named. You may review all, some or none of these projects, in any order. Each project review usually takes about 5–10 minutes. The first review may take longer while you become familiar with the classifications.

For each project, you will be shown the public register entry and the classifications proposed by the system:

- one or more **Research Domains**, describing what the project is about;
- one or two **Analytical Purposes**, describing what the project is trying to do analytically; and
- two binary **cross-cutting tags**—Demographic disparities / equity and COVID-19 & Pandemic—each shown as either Applied or Not applied.

You will be asked whether each proposed classification accurately describes the project and whether the basis for it is visible in the public project title and datasets listed in the register. You can also identify any important Domains, Purposes or tags that are missing, and comment on whether the public register entry contains enough information and whether the available classification categories represent the project adequately.

Most free-text comments are optional. A brief explanation is required only when you say that a proposed Domain or Purpose does not fit, or that a proposed tag status is incorrect. Please provide only general project context and do not disclose confidential, sensitive or otherwise non-public information.

You may save an unfinished review using **Save & Return Later** and return through the same personalised link. Please do not forward the link.

- **What will my responses be used for?**

Your responses are analysed alongside those of other project researchers and of trained independent coders, who classify the same public register entries without access to your answers. Together this evidence is used to assess how well the proposed classifications fit the projects, whether the public register entry contains enough information to classify reliably, and whether the classification framework covers projects adequately.

Your review is treated as project-informed evidence rather than as a definitive 'correct' answer, because you may know things about your project that are not visible in the public register entry. How findings are reported, and whether you could be identified, is explained below.

Is participation voluntary?

Yes. Taking part is entirely voluntary. Deciding not to participate will have no adverse consequences. You may skip any project review or stop at any time.

Are there any risks?

This is a low-risk professional review. Some questions may involve mild discomfort when commenting on the accuracy of a public project description or classification. Your review is not an assessment of you or of the quality of your research. You should not disclose confidential, sensitive or otherwise non-public information. Where wider project context affects your answer, please describe it only at a general level you are comfortable sharing. Optional comments may be left blank. You should be aware that the group invited to this study is small, and that all participants are publicly named on publicly listed projects. This means your responses cannot be treated as fully anonymous. A reader familiar with the register might be able to infer who commented on a particular project. We describe below the steps we take to limit this, and the choices available to you.

What personal data do you collect about me?

We collect and process the following:

- your name and professional email address, obtained from publicly available sources;
- a record of your consent, including the date, time and the version of this information sheet you were shown;
- your survey responses, including any optional free-text comments;
- a pseudonymous participant identifier and the key linking it to your contact details;
- your optional preference about acknowledgement.

How is my information kept confidential?

Your name and contact details are held in a restricted recruitment table, stored separately from your survey responses. Your responses are recorded against a pseudonymous identifier, not your name. Analysis datasets contain no direct identifiers. Access to identifiable contact and consent data is restricted to Dr Balint Stewart for study administration; Dr Balint Stewart and Dr Joseph Lam may access pseudonymised research data. Response data are held in UCL REDCap and other UCL-approved secure systems, protected by password-controlled accounts.

Will I be identified in any outputs?

Findings will be reported in aggregate. Neither your name nor your responses will be published in identifiable form without your specific permission.

Occasionally we may describe an anonymised example to illustrate a type of problem. Any such description focuses on the classification issue rather than the project, and omits the project title, the datasets used, and any other detail we judge would make the project recognisable. Your own comments are not reproduced.

If we would like to use a short quotation from your comments, we will contact you by email with the exact proposed wording and the context in which it would appear, and we will use it only if you agree in writing. You may decline without giving a reason and without affecting anything else about your participation.

Because participants are publicly named on publicly listed projects, we cannot guarantee that a quotation would be entirely untraceable to you by a reader who knows the register well. Sending you the exact wording in advance is intended to let you judge that for yourself before deciding.

Separately, you may choose whether you would like to be acknowledged by name in the resulting publications. This choice is entirely optional, made independently of your decision to take part, and does not affect your participation. Please note that if you choose to be acknowledged, your involvement in this study becomes permanently and publicly identifiable, and this cannot be reversed once material is published. If you would prefer not to be named, simply decline the acknowledgment and the study team will not name or acknowledge you in resulting outputs. Because participants are publicly associated with listed projects, complete anonymity cannot be guaranteed.

How long will my data be kept?

Your contact details and the linking key will be securely deleted within three months of the study ending, once withdrawal requests and participant communications are complete. Consent records will be retained only as long as required for research governance.

Pseudonymised research data will be retained securely for 10 years so that findings can be verified and, where appropriate, re-analysed by the research team. Any wider sharing or archiving will be of anonymised aggregate data only; no identifiable data will be shared. Your individual responses will not be shared with researchers outside the study team, deposited in a public repository, or shared with any commercial organisation.

Can I withdraw?

Yes. You may stop at any time, skip any project review, or ask not to receive further invitations. Any review you have not submitted will not be analysed.

You may also request withdrawal of a review you have already submitted. To do this, email balint.stewart@ucl.ac.uk quoting the Review reference shown at the top of that review. You may request one review to be withdrawn, or all of them.

Requests must be received by Monday 19 October 2026, which is 14 calendar days after the survey closes on Monday 5 October. After this date, the analysis dataset will be locked and responses may already have been incorporated into the analysis, so individual reviews will no longer be removed

Will I hear about the results?

Yes. All participants will be offered a brief plain-language summary of the overall findings by email once the study is complete. Findings may also be shared through the public dashboard, in academic publications and conference presentations, on the Open Science Framework, and through UCL or funder communications. Results will be reported in aggregate.

Who can I contact?

If you have questions about the study, please contact:

Dr Balint Stewart (study contact), balint.stewart@ucl.ac.uk

Dr Joseph Lam (Principal Investigator), joseph.lam@ucl.ac.uk

This research is funded by Health Data Research UK and is being conducted at University College London.

Complaints. If you have a concern about any aspect of this study, please contact Dr Balint Stewart in the first instance, or Dr Joseph Lam, the Principal Investigator, at joseph.lam@ucl.ac.uk. If you feel your complaint has

not been handled to your satisfaction, you can contact the Chair of the UCL Research Ethics Committee at ethics@ucl.ac.uk.

Data protection queries. UCL's Data Protection Officer can be contacted at data-protection@ucl.ac.uk.

Local Data Protection Privacy Notice

The controller for this project will be University College London (UCL). The UCL Data Protection Officer provides oversight of UCL activities involving the processing of personal data, and can be contacted at data-protection@ucl.ac.uk.

This 'local' privacy notice sets out the information that applies to this particular study. Further information on how UCL uses participant information can be found in our 'general' privacy notice for participants in research studies: <https://www.ucl.ac.uk/legal-services/privacy/ucl-general-research-participant-privacy-notice>

The information that is required to be provided to participants under data protection legislation (GDPR and DPA 2018) is provided across both the 'local' and 'general' privacy notices.

The categories of personal data used will be as follows: name; professional email address; record of consent, including date, time and information sheet version; survey responses, including optional free-text comments; a pseudonymous participant identifier and the key linking it to contact details; and your optional preference about acknowledgement.

The lawful basis that will be used to process your personal data is performance of a task in the public interest.

No special category personal data will be collected in this study.

Your personal data will be processed only so long as it is required for the research project. Contact details and the linking key will be deleted within three months of the study ending. Responses are held against a pseudonymous identifier, and we minimise the processing of personal data wherever possible.

Your personal data will not be transferred outside the UK, and will not be shared with any recipient outside the UCL research team.

If you are concerned about how your personal data is being processed, or if you would like to contact us about your rights, please contact UCL in the first instance at data-protection@ucl.ac.uk.